

Homework 4: Person Name ('Named Entity') Classification

[starter code]

Due: Tuesday April 10 11:59pm Pacific Daylight Time

For this assignment, you will be building a maximum entropy Markov model (MEMM) for identifying person names in newswire text. We have provided all of the machinery for training and testing your MEMM, but we have left the feature set woefully inadequate. Your job is to modify the code for generating features so that it produces a much more sensible, complete, and higher-performing set of features.

If you're programming in Python, please use Python 2.6 or higher because the program uses JSON module that was introduced in version 2.6. Also, you're required to have Java working on your system because a significant amount of the starter codes are written in Java. However, you do not need to worry about the Java part because we provide the commands to compile and run the program.

Starter Code

In the starter code we have provided you with three simple starter features, but you should be able to improve substantially on them. We recommend experimenting with orthographic information, gazeteers, and the surrounding words, and we also encourage you to think beyond these suggestions.

For Java, the file you will be modifying is `FeatureFactory.java`. The class currently looks something like this:

```
public class FeatureFactory {

    private static List<String> computeFeatures(List<String> words, String previousLabel, int position) {

        List<String> features = new ArrayList<String>();

        String currentWord = words.get(position);

        // Baseline Features
        features.add("word=" + currentWord);
        features.add("prevLabel=" + previousLabel);
        features.add("word=" + currentWord + ", prevLabel=" + previousLabel);
```

```
// TODO: Add your features here

return features;
}
}
```

For Python, the file you will be modifying is `FeatureFactory.py`.

```
def computeFeatures(self, words, previousLabel, position):
    features = []
    currentWord = words[position]

    """ Baseline Features """
    features.append("word=" + currentWord)
    features.append("prevLabel=" + previousLabel)
    features.append("word=" + currentWord + ", prevLabel=" + previousLabel)

    """ TODO: Add your features here """

    return features
```

Adding Features to the Code

You will create the features for the word at the given position, with the given previous label. You may condition on any word in the sequence (and its relative position), not just the current word, because they are all observed. You may not condition on any labels other than the previous one. Each function you build will be a binary function of some kind of feature. The features are stored in a list because we are using a *sparse* representation. Features which have a value of *true* (or `1.0`) will be present in the list and everything not present is assumed to be false (or `0.0`).

The argument to `features.add()` in Java or `features.append()` in Python is the name of the feature. You need to give a unique name for each feature. The system will use this unique name in training to set the weight for that feature. At testing time the system will use the name of this feature and its weight to make a classification decision.

Consider the following sample feature:

```
features.add("word="+currentWord);
```

and imagine that the current word is "Oakland".

This means "I am a binary function that is true if the current word is 'Oakland'". For every such feature you write, the code will take this function and actually learn two separate weights for this feature: one for "label=PERSON" and one for "label=O".

So just by specifying this feature, the system will learn two functions, call them `f1` and `f2`, where `f1`, which means "I am a binary function that is true if the current word is 'Oakland' and the output label is

'Person'" will get a low weight, and f2, which means "I am a binary function that is true if the current word is 'Oakland' and the output label is 'O'" will get a high weight. You can think of the name of both f1 and f2 as the string "word=Oakland" and not worry about the fact that we are learning two separate weights. Then, when the system is running on a test sentence that has the word Oakland, it will call your code to create the feature "word=Oakland", and it will automatically create two features, "word=Oakland,label=PERSON" and "word=Oakland,label=O", and it will use those weights in the classifier.

Similarly, this sample feature:

```
features.add("word=Jenny, prevLabel=O");
```

will create two features, each with its own weight, one for "word=Jenny, prevLabel=O, label=PERSON" and one for "word=Jenny, prevLabel=O, label=O".

Types of features to include

Your features should not just be the words themselves. The features can represent any property of the word, context, or additional knowledge.

For example, the case of a word is a good predictor for a person's name, so you might want to add a feature to capture whether a given word was lowercase, Titlecase, CamelCase, ALLCAP, etc.

Imagine you saw the word 'Jenny'. In addition to the feature for the word itself (as above), you could add a feature to indicate it was in Title case, like:

for Java,

```
if (Character.isUpperCase(currentWord.charAt(0))) {  
    features.add("case=Title");  
}
```

for Python,

```
if currentWord[0].isupper():  
    features.append("case=Title")
```

It will create two features, each with its own weight, one for "case=Title, label=PERSON" and one for "case=Title, label=O". It is easy to see why this helps. You might encounter an unknown word in the test set, but if you know it begins with a capital letter then this might be evidence that helps with the correct prediction.

Choosing the correct features is an important part of natural language processing. It is as much art as science: some trial and error is inevitable, but you should see your accuracy increasing as you add new types of features.

The name of a feature is not different from an ID number. You can use assign any name for a feature as long as it is unique. For example, you can use "Titlecase" instead of "case=Title".

Running the Program

For Java, execute

```
$ cd java
$ mkdir classes
$ javac *.java org/json/*.java -d classes -target 1.6
$ java -cp classes -Xmx1G NER ../data/train ../data/dev
```

Similarly for Python, execute

```
$ cd python
$ mkdir classes
$ javac *.java org/json/*.java -d classes
$ python NER.py ../data/train ../data/dev
```

We have provided you with a training set, called `train` and a development test set called `dev`. We will be running your programs on an unseen test set, so you should try to make your features as general as possible. Your goal should be to increase F1, which is the harmonic mean of the precision and the recall. If you run the program as-is, you should get the following score:

```
precision = 0.802
recall = 0.522
F1 = 0.633
```

When you run the program, you will see it print out a lot of information as it does the optimization, and you can pretty much ignore that. Afterwards, it will print out your score as above. You can give it an additional flag, `-print` and have it print the test set along with the real answers and your guesses. The first column is the word, the second column is the true answer, and the third column is your program's guess. This should help you do error analysis to see what kinds of things you are getting right, and what kind you are getting wrong. This will help you properly target your features.

For Java, execute

```
$ java -cp classes -Xmx1G NER ../data/train ../data/dev -print
```

For Python, execute

```
$ python NER.py ../data/train ../data/dev -print
```

Submission

For Java, execute

```
$ cd java
$ mkdir classes
$ javac *.java org/json/*.java -d classes
$ java -cp classes -Xmx1G Submit
```

For Python, execute

```
$ cd python
$ mkdir classes
$ javac *.java org/json/*.java -d classes
$ python submit.py
```

When you encounter "Out of Heap Space" problem in the attempt to run or submit the program in Java, increase the heap size by using "-Xmx2G" instead of "-Xmx1G" option.

It is important that these scripts be executed from either the java or python directories so that they can easily import the code you've written and so that the relative path "../data/" correctly points to the training data. It is hard-coded into the submission scripts. You are allowed to submit as many times as you want. We will take the max of your scores. The submission takes twice longer than the usual running time of your program since your program will be run for both dev and test data sets. Please make sure to allow enough time for the submit script to run.

Grading

The assignment will be graded based on its F-scores on both development and held-out data sets. 10 points will be given to the performance of each data set. For the development data set, achieving 70% F-score would be fairly easy and 80% would require a reasonable amount of work. If you're achieving above 89%, it's very impressive. Please use these numbers as a ball park and be careful not to over-fit your features to the development data set.

Tips

- Start early. This assignment may take longer than the previous assignments if you're aiming for the perfect score.
- Generalize your features. For example, if you're adding the above "case=Title" feature, think about whether there is any pattern that is not captured by the feature. Would the "case=Title" feature capture "O'Gorman"?
- When you add a new feature, think about whether it would have a positive or negative weight for PERSON and O tags (these are the only tags for this assignment).
- "-print" option is useful when you want compare the differences of the predictions before and after adding new features.